Module 3, Lesson 2: Mongoid Relationships

The overall goal of the assignment is to give you practice in:

- Defining models, custom type classes, and relationships
- Defining relationship cardinality (one and many), realization (embedded and linked), and navigation types (uni and bi-directional)
- Forming and manipulating relationships

The functional goal of the assignment is to:

- Implement a data tier for managing `Contest` information to include the following document types:
  - `Venue`, `Judge`, `MedicalRecord`, and `Racer`

Note that this assignment was written so that you can implement it in parts after each lecture. If you are performing the assignment in between lectures, stop at the next lecture boundary in the technical requirements section and resume once you have completed the lecture. You are free to experiment with other forms of the configurations presented, but the grading will only be targeted at the specific requirements listed.
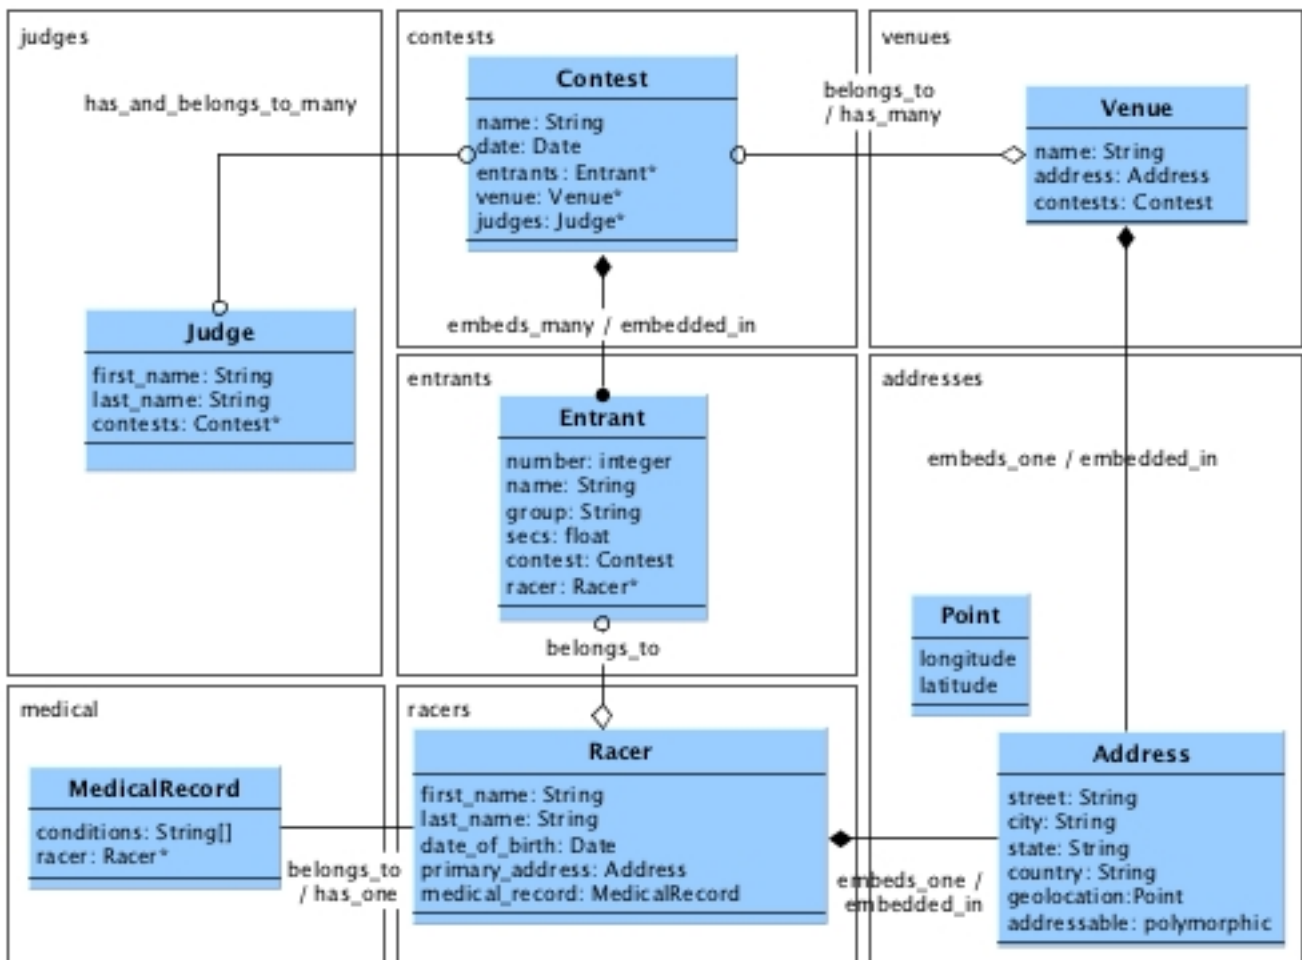
Functional Requirements

Figure 1: Model/Relationships

- **Diagram Notes**:
  - Model/Relationships Image
  - (*) represents a foreign key

– the upper left-hand names in each boundary box represents the corresponding collection names within MongoDB

1. Implement a (custom type) `Point` that encapsulates the geographic coordinates of an `Address`.

2. Implement a (1:1 embedded) relationship between `Racer`/`Venue` and `Address`.

3. Implement a (1:M linked) relationship between `Venue` and `Contest` and an embedded linked relationship between `Entrant` and `Racer`.

4. Implement a (1:M embedded) relationship between `Contest` and `Entrant`.

5. Implement a (1:1 Linked) relationship between `Racer` and `MedicalRecord`.

6. Implement a (M:M Linked) relationship between `Judge` and `Contest`.

Getting Started

1. Start your MongoDB server using `mongod`

2. Create a new Rails application called `contests`.

   ```
   $ rails new contests
   $ cd contests
   ```

3. Setup your application for Mongoid.

   - Add the `mongoid`, `rspec-rails`, and `mongoid-rspec` gems to your Gemfile and run `bundle`. Notice that the `:test` group encapsulates the two related `rspec` gems.

     ```
     gem 'mongoid', '~> 5.0.0'

     group :test do
         gem 'rspec-rails', '~> 3.0'
         gem 'mongoid-rspec', '3.0.0'
     end
     ```

   - Generate a `mongoid.yml` configuration file

     ```
     $ rails g mongoid:config
         create config/mongoid.yml
     ```

   - Add the generated `mongoid.yml` file to `config/application.rb`

     ```
     module Contests
         class Application < Rails::Application
         ...
         # Boostraps mongoid within applications -- like rails console
         Mongoid.load!('./config/mongoid.yml')
     ```

4. Download and extract the starter set of files. The root directory of this starter set will be referred to as the root directory of your solution. When extracted correctly – the `spec` folder should be at the same (root) level.

   ```
   --- student-start
       |-- .rspec (important hidden file)
       `-- spec
           |-- test_utils.rb
           |-- customtype_spec.rb
           |-- lecture1_spec.rb
           |-- lecture2_spec.rb
           |-- lecture3_spec.rb
           |-- lecture4_spec.rb
           |-- lecture5_spec.rb
           |-- rails_helper.rb
           `-- spec_helper.rb
   ```

- spec - this directory contains tests to verify your solution. You should not modify anything in this directory

5. Implement the technical requirements.

6. Run the rspec command from the project root directory. The spec files are written per-lecture. The steps taken in one lecture can impact the results of a preceding lecture. However, if you execute all sections correctly, you will be able to execute all rspec tests at the end and pass.

```
$ rspec

(N) examples, (N) failures
...
```

Technical Requirements

Setup: Custom Type

In this section we will perform some catch-up and implement a custom type described in Lesson 1. Custom types have structure but no `_id`. They are used as a convenience for handling fields and marshaling/de-marshaling those fields. Otherwise your application code can work directly with hashes.

1. Add a custom type called `Point` (in the `app/model` directory) that represents the longitude and latitude coordinates for an `Address` expressed in `GeoJSON` format. This class must:

   - be called `Point`
   - have a read/write attribute called `longitude`
   - have a read/write attribute called `latitude`
   - have an initializer that accepts the two attribute values in the order of `longitude` and `latitude`

   You can use the rails console to demonstrate your new class and initializer method.

   ```
   > Point.new(0,1)
    => #<Point:0x000000053e69d0 @longitude=0, @latitude=1>
   ```

   ```
   $ rspec spec/customtype_spec.rb -e rq01
   ```

2. Add an instance method to the `Point` class called `mongoize` that returns a hash in `GeoJSON Point` format and ready to be stored within MongoDB. An instance method by this name is required by Mongoid for use of custom types. This method must:

   - accept no arguments
   - return a hash with a
     - key `type` with the String value `"Point"` and
     - key `coordinates` with an array containing `longitude` and `latitude` in that order.

   You can use the rails console to demonstrate your new instance method.

   ```
   > Point.new(0,1).mongoize
    => {:type=>"Point", :coordinates=>[0, 1]}
   ```

   ```
   $ rspec spec/customtype_spec.rb -e rq02
   ```

3. Add a class method called `demongoize` that will return an instance of the `Point` class initialized from the contents in the database. A class method by this name is required by Mongoid for custom types. This method must:

   - accept a hash object assumed to be the result of the `mongoize` method
   - extract the `longitude` and `latitude` from the `coordinates` array in the hash
   - instantiate a new Point instance with the `longitude` and `latitude`
   - return the new Point instance

   You can use the rails console to demonstrate your new class method.

   ```
   > Point.demongoize(:type=>"Point", :coordinates=>[0, 1])
    => #<Point:0x0000000538e7a8 @longitude=0, @latitude=1>
   ```

   ```
   $ rspec spec/customtype_spec.rb -e rq03
   ```

3

4. Add a class method called `mongoize` that will accept either `Point` or hash types and return a `mongoized` string for either inputs. A class method by this name is required by Mongoid for custom types. This method must:

   - accept a single input
   - determine the type of the input
   - create a `mongoized` form of the `Point` ready to be store in MongoDB.

   You can use the rails console to demonstrate your new class method.

   ```
   > Point.mongoize(Point.new(0,1))
    => {:type=>"Point", :coordinates=>[0, 1]}
   > Point.mongoize(:type=>"Point", :coordinates=>[0, 1])
    => {:type=>"Point", :coordinates=>[0, 1]}
   ```

   ```
   $ rspec spec/customtype_spec.rb -e rq04
   ```

5. Add a class method called `evolve` that performs the same functionality as `mongoize`. A class method by this name is required by Mongoid for custom types.

   You can use the rails console to demonstrate your new class method.

   ```
   > Point.evolve(Point.new(0,1))
    => {:type=>"Point", :coordinates=>[0, 1]}
   > Point.evolve(:type=>"Point", :coordinates=>[0, 1])
    => {:type=>"Point", :coordinates=>[0, 1]}
   ```

   ```
   $ rspec spec/customtype_spec.rb -e rq05
   ```

6. Use the `rails generate` to create a model class called `Address` with the following fields:

   - `street` : String
   - `city` : String
   - `state` : String
   - `country` : String
   - `geolocation` : Point

   You can use the rails console to demonstrate your new model class with the embedded custom type.

   ```
   > Address.create(:geolocation=>Point.new(0,1))
    => #<Address _id: 5675f08de301d0a1fb000000,
       street: nil, city: nil, state: nil, country: nil,
       geolocation: {:type=>"Point", :coordinates=>[0, 1]}>

   > Address.collection.find.first
    => {"_id"=>BSON::ObjectId('5675f08de301d0a1f    Can navigate from Racer to Address (FAILED - 2)
   b000000'),
       "geolocation"=>{"type"=>"Point", "coordinates"=>[0, 1]}}

   > Address.first.geolocation
    => #<Point:0x00000047aee10 @longitude=0, @latitude=1>
   ```

   ```
   $ rspec spec/customtype_spec.rb
   ```

Lecture 1: 1:1 Embedded

In this section we will build a 1:1 embedded relationship between `Racer`/`Venue` and `Address`. We will start with just `Racer` and `Address` and build a concrete 1:1 embedded relationship. It will eventually evolve to a 1:1 embedded polymorphic relationship since Address is not specific to `Racer`.

1. Use the `rails generate` to create a model class called `Racer` with the following fields:

   - `first_name` : String – mapped to the document field `fn`
   - `last_name` : String – mapped to the document field `ln`
   - `date_of_birth` : Date – mapped to the document field `dob`

4

You can use the rails console to demonstrate your new model class, fields, and field mappings.

```
> Racer.new
 => #<Racer _id: 5675fb8ae301d0a1fb000002,
    first_name(fn): nil, last_name(ln): nil, date_of_birth(dob): nil>

$ rspec spec/lecture1_spec.rb -e rq01
```

2. Define a 1:1 uni-directional, embeded relationship from `Racer` to `Address` using the `embeds_one` macro. (**Hint**: Since the relationship name `primary_address` is different from the classname `Address`, you must specify a `class_name` mapping). This relationship must:

   - be called `primary_address`
   - embed the `Address` instance within the body of the `Racer`

You can use the rails console to demonstrate your new relationship. In the example below, we are forming the relationship at the point where query is executed.

```
> address=Address.new(:city=>"somewhere", :geolocation=>Point.new(0,1))
 => #<Address _id: 5675fedfe301d0a1fb000003, street: nil, city: "somewhere",
    state: nil, country: nil, geolocation: {:type=>"Point", :coordinates=>[0, 1]}>
> r=Racer.create(:fn=>"cat",:ln=>"inhat",:primary_address=>address)
 => #<Racer _id: 5675ff11e301d0a1fb000004,
    first_name(fn): "cat", last_name(ln): "inhat", date_of_birth(dob): nil>
```

Note the `Address` instance is embedded within the `Racer` instance.

```
> pp r.attributes
{"_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004'),
 "first_name"=>"cat",
 "last_name"=>"inhat",
 "primary_address"=>
  {"_id"=>BSON::ObjectId('5675fedfe301d0a1fb000003'),
   "city"=>"somewhere",
   "geolocation"=>{:type=>"Point", :coordinates=>[0, 1]}}}
```

Notice that you can navigate from the `Racer` to the `Address` fields. Since this is currently uni-directional, we cannot navigate from the `Address` instance to the containing `Racer` instance.

```
> r.first_name
 => "cat"
> r.primary_address.city
 => "somewhere"
> r.primary_address.geolocation
 => #<Point:0x00000004ffe818 @longitude=0, @latitude=1>

$ rspec spec/lecture1_spec.rb -e rq02
```

3. Define a 1:1 concrete relationship in the reverse direction from `Address` to `Racer` using the `embedded_in` macro. This relationship must:

   - be called `racer` and be specific to `Racer` (i.e., we are not yet worried about Venue at this point)

You can demonstrate your new bi-directional relationship by navigating back and forth between the `Racer` and `Address`.

```
> r=Racer.where(:fn=>"cat",:ln=>"inhat").first
 => #<Racer _id: 5675ff11e301d0a1fb000004,
    first_name(fn): "cat", last_name(ln): "inhat", date_of_birth(dob): nil>
> address=r.primary_address
> address.racer
 => #<Racer _id: 5675ff11e301d0a1fb000004,
    first_name(fn): "cat", last_name(ln): "inhat", date_of_birth(dob): nil>

$ rspec spec/lecture1_spec.rb -e rq03
```

4. Generalize the 1:1 concrete relationship using a polymorphic construct such that any model class can also embed `Address` instance(s). The `Address` class must

- change the name of the relationship from `racer` to `addressable`
- annotate the relationship as being `polymorphic`

The `Racer` class must:

- annotate the relationship on its end as `addressable`

You may use the rails console to repeat the actions of the previous step to demonstrate your new polymorphic relationship.

```
> r=Racer.where(:fn=>"cat",:ln=>"inhat").first
> address=r.primary_address
> address.addressable
=> #<Racer _id: 5675ff11e301d0a1fb000004,
   first_name(fn): "cat", last_name(ln): "inhat", date_of_birth(dob): nil>

$ rspec spec/lecture1_spec.rb -e rq04
```

5. Use the `rails generate` command to create a model class called `Venue` with the following fields:

- `name` : String

You can use the rails console to demonstrate your new model class.

```
> Venue.new
=> #<Venue _id: 56760a1ce301d0a1fb000005, name: nil>

$ rspec spec/lecture1_spec.rb -e rq05
```

6. Create a 1:1 embedded relationship between `Venue` and `Address` using the now polymorphic `addressable` relationship. This relationship must:

- be called `address`
- embed an instance of `Address` within `Venue` instances
- annotate the relationship as `addressable`

You can demonstrate your new relationship using the rails console. If you inspect the `Venue` class for methods containing the word `address` (**Hint**: `Venue.methods.grep /address/`) you can locate alternatives for creating the relationship.

```
> v=Venue.create(:name=>"Boston")
=> #<Venue _id: 56760ddbe301d0a1fb00000a, name: "Boston">
> Venue.find_by(:name=>"Boston").attributes
=> {"_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a'), "name"=>"Boston"}

> v.create_address(:city=>"Boston", :state=>"MA", :geolocation=>Point.new(71.5,42.21))
=> #<Address _id: 56760de9e301d0a1fb00000b, street: nil, city: "Boston", state: "MA",
   country: nil, geolocation: {:type=>"Point", :coordinates=>[71.5, 42.21]}>
> pp Venue.find_by(:name=>"Boston").attributes
{"_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a'),
 "name"=>"Boston",
 "address"=>
  {"_id"=>BSON::ObjectId('56760de9e301d0a1fb00000b'),
   "city"=>"Boston",
   "state"=>"MA",
   "geolocation"=>{"type"=>"Point", "coordinates"=>[71.5, 42.21]}}}

> v.address.addressable
=> #<Venue _id: 56760ddbe301d0a1fb00000a, name: "Boston">
```

Note that setting the `address` to `nil` removes the `Address` from the `Venue` document.

```
> v.address=nil
 => nil
> pp Venue.find_by(:name=>"Boston").attributes
{"_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a'), "name"=>"Boston"}

$ rspec spec/lecture1_spec.rb -e rq06
```

Lecture 2: 1:M Linked

In this section you will be asked to create a few examples of 1:M linked relationships. The first one between `Contest` and `Venue` will be straight forward. The second one between `Entrant` and `Racer` will be complicated by `Entrant` becoming an embedded type in the next section.

1. Use the `rails generate` command to create a model class called `Contest` with the following fields:

   - `name` : String
   - `date` : Date

   Also add:

   - a mixin to track updates, but not creates (**Hint**: `Mongoid::Timestamps::Updated`)

   You can use the rails console to demonstrate your new model class, fields, and field mappings.

   ```
   > Contest.new
    => #<Contest _id: 56761701e301d0a1fb00000c, updated_at: nil, name: nil, date: nil>

   $ rspec spec/lecture2_spec.rb -e rq01
   ```

2. Create a M:1 linked, uni-directional relationship from `Contest` to `Venue` using the `belongs_to` macro. This relationship must:

   - be called `venue`
   - form a foreign key link from the `Contest` to the `Venue`

   You can use the rails console to demonstrate your new M:1 uni-directional relationship.

   ```
   > v=Venue.find_by(:name=>"Boston")
    => #<Venue _id: 56760ddbe301d0a1fb00000a, name: "Boston">
   > c=Contest.create(:name=>"Boston 5K", :date=>Date.new(2015,5,30))
   > pp Contest.find_by(:name=>"Boston 5K").attributes
   {"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
    "name"=>"Boston 5K",
    "date"=>2015-05-30 00:00:00 UTC,
    "updated_at"=>2015-12-20 02:58:12 UTC}
   ```

   Notice that when you asssign the `Venue` to the `Contest` and save the `Contest`, the foreign key to the `Venue` is written into the `Contest` instance.

   ```
   > c.venue = v
    => #<Venue _id: 56760ddbe301d0a1fb00000a, name: "Boston">
   > c.save
   > pp Contest.find_by(:name=>"Boston 5K").attributes
   {"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
    "name"=>"Boston 5K",
    "date"=>2015-05-30 00:00:00 UTC,
    "updated_at"=>2015-12-20 03:02:00 UTC,
    "venue_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a')}
   ```

   Notice that if you reset the state of the in-memory instances and access the `Contest` again, the `Venue` is not accessed.

   ```
   > reload!

   > c=Contest.find_by(:name=>"Boston 5K")
   ```

```
D, | {"find"=>"contests", "filter"=>{"name"=>"Boston 5K"}}
 => #<Contest _id: 56761944e301d0a1fb00000d, updated_arailst: 2015-12-20 03:02:00 UTC, name: "Boston 5K
```

We can even access the `_id` of the `Venue` from the `Contest` without loading the `Venue`. This is because Mongoid gives us direct access to the foreign key property.

```
> c.venue_id
 => BSON::ObjectId('56760ddbe301d0a1fb00000a')
```

However, if we cross the relationship – the `Venue` is loaded.

```
> c.venue.id
D, | {"find"=>"venues", "filter"=>{"_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a')}}
 => BSON::ObjectId('56760ddbe301d0a1fb00000a')
```

```
$ rspec spec/lecture2_spec.rb -e rq02
```

3. Implement the inverse side of the M:1 relationship from `Venue` to `Contest` using the `has_many` macro. This does not add a foreign key local to `Venue`. It causes `Venue` to search for `Contests` with its primary key listed as a foreign key. The relationship must:

   - be called `contests`

You can demonstrate your new 1:M inverse relationship using the rails console. Pay close attention to the debug printed (and not printed) by the driver. We start by getting a `Contest` with an association with a `Venue` to obtain the primary key for the `Venue`.

```
> reload!
> c=Contest.find_by(:name=>"Boston 5K")
> vid=c.venue_id
```

When we access the `Venue` via the find method, the `venues` collection is searched for documents with a primary key (`_id`) equal to our value.

```
> v=Venue.find(vid)
D, | {"find"=>"venues", "filter"=>{"_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a')}}
 => #<Venue _id: 56760ddbe301d0a1fb00000a, name: "Boston">
```

When we access the `Contests` via the `venue.contests` method, the `contests` collection is searched for foreign keys (`venue_id`) equal to our value.

```
> v.contests.map {|c| c.name}
D, | {"find"=>"contests", "filter"=>{"venue_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a')}}
 => ["Boston 5K"]
```

Prior to removing the relationship, we can check the state of the database and see the foreign key in the `Contest` collection.

```
> pp c.attributes
{"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
 "name"=>"Boston 5K",
 "date"=>2015-05-30 00:00:00 UTC,
 "updated_at"=>2015-12-20 03:02:00 UTC,
 "venue_id"=>BSON::ObjectId('56760ddbe301d0a1fb00000a')}
```

When we remove the `Contest` (many-side) from the `Venue` (one/inverse-side) side of the bi-directional relationship, the foreign key in the `Contest` is updated along with the `updated_at` timestamp.

```
> v.contests.delete c
D, | {"update"=>"contests", "updates"=>[{"q"=>{"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d')},
   "u"=>{"$set"=>{"venue_id"=>nil, "updated_at"=>2015-12-20 03:29:54 UTC}}...
```

We are left with a `nil` foreign key in the `Contest` with an updated `updated_at` timestamp.

```
> pp Contest.find(c.id).attributes
{"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
 "name"=>"Boston 5K",
```

```
  "date"=>2015-05-30 00:00:00 UTC,
  "updated_at"=>2015-12-20 03:29:54 UTC,
  "venue_id"=>nil}
```

Further accesses for `Contest` instances for the `Venue` come up empty.

```
> v.contests.map {|c| c.name}
 => []
```

```
$ rspec spec/lecture2_spec.rb -e rq03
```

4. Use the `rails generate` command to create a model class called `Entrant` with the following fields:

   - `_id` : Integer – assigned to hold the `number` of the entrant
   - `name` : String – cached name of the `Racer`
   - `group` : String
   - `secs` : Float

You can demonstrate your new model class using the rails console.

```
> Entrant.new(:id=>1)
 => #<Entrant _id: 1, name: nil, group: nil, secs: nil>
```

```
$ rspec spec/lecture2_spec.rb -e rq04
```

5. Add a M:1 uni-directional, linked relationship from `Entrant` to `Racer` using the `belongs_to` macro. This relationship must:

   - be called `racer`
   - copy the name of the `Racer` (by concatenating "last, first" names separated by a comma, into a single string) when the instance is created. (Hint: `before_create` callback).

You can demonstrate your new relationship using the rails console. When we associate a new `Entrant` with a `Racer` instance, only the primary key of the `Racer` is stored.

```
> racer=Racer.find_by(:fn=>"cat",:ln=>"inhat")
> entrant=Entrant.new(:id=>1, :racer=>racer, :group=>"masters")
 => #<Entrant _id: 1, name: nil, group: "masters", secs: nil,
    racer_id: BSON::ObjectId('5675ff11e301d0a1fb000004')>
```

However, when the `Entrant` is saved, the `before_create` callback is invoked, where the `Entrant` has a chance to copy off the `Racer`'s first/last name into a field within the `Entrant` document. It is assumed that `Entrants` never change `Racer` instances and `Racer` names rarely change. However, it is also assumed that the `Racer` name will be needed in `Contest` race results.

```
> entrant.save
D, | {"insert"=>"entrants", "documents"=>[{"_id"=>1, "group"=>"masters",
   "racer_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004'), "name"=>"inhat, cat"}]
```

Note that the name of the `Racer` can be obtained from the cached copy within the `Entrant`.

```
> entrant.name
 => "inhat, cat"
```

```
$ rspec spec/lecture2_spec.rb -e rq05
```

6. Implement the inverse side of the M:1 linked relationship from `Racer` to `Entrant` using the `has_many` macro. This will be a temporary construct because in the next section we will be changing `Entrant` to an embedded class and will need to do something different. For now, this relationship must:

   - be called `races` (**Hint**: this relationship must also be mapped to the `Entrant` class since the name `Entrant` and the singular form of `races` is not the same.

You can demonstrate your new bi-directional relationship using the rails console. Notice if more `Entrant` associations are made with `Racer` – they are all located using a single foreign key query.

```
> (2..3).each {|index| Entrant.create(:id=>index, :racer=>racer)}
> racer.races.map {|r| r.id}
```

```
D, | {"find"=>"entrants",
  "filter"=>{"racer_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004')}}
 => [2, 3]
```

```
$ rspec spec/lecture2_spec.rb -e rq06
```

Lecture 3: 1:M Embedded

In this section you will focus on implementing a 1:M embedded relationship. In this case, the embedded document will be an "annotated link" from `Contest` to `Racer` augmented with racer-specific information for the contest using an `Entrant` class.

1. Create an 1:M, uni-directional embedded relationship from `Contest` to `Entrant` using the `embeds_many` macro. This relationship must:

   - be called `entrants`

   You can demonstrate your new relationship using the rails console. However, until we make this a bi-directional relationship there will not be much you can do with the relationship.

   ```
   > Contest.new.entrants
    => []
   ```

   ```
   $ rspec spec/lecture3_spec.rb -e rq01
   ```

2. Complete the M:1, bi-directional embedded relationship from `Entrant` to `Contest` using the `embedded_in` macro. This relationship must:

   - be called `contest`

   You can demonstrate your new relationship using the rails console.

   ```
   > Entrant.new.embedded?
    => true
   > Entrant.new.contest
    => nil
   ```

   However, we soon will see that embedding `Entrant` within `Contest` causes a problem with the inverse side of the 1:M relationship from `Racer` to 'Entrant'.

   ```
   > racer=Racer.find_by(:fn=>"cat",:ln=>"inhat")
   > racer.races
     Mongoid::Errors::MixedRelations:
     message:
       Referencing a(n) Entrant document from the Racer document via
       a relational association is not allowed since the Entrant is
       embedded.
     summary:
       In order to properly access a(n) Entrant from Racer the reference
       would need to go through the root document of Entrant. In a simple
       case this would require Mongoid to store an extra foreign key
       for the root, in more complex cases where Entrant is multiple
       levels deep a key would need to be stored for each parent up
       the hierarchy.
     resolution:
       Consider not embedding Entrant, or do the key storage and access
       in a custom manner in the application code.
   ```

   Note that many of the previous rspec tests that succeeded will now fail as the result of the current state detailed above. This will be repaired in the next step.

   ```
   $ rspec spec/lecture3_spec.rb -e rq02
   ```

3. Modify the M:1 relationship from `Entrant` to `Racer` to be uni-directional again, thus only directly navigatable from the embedded `Entrant` owning side. We will provide an alternative once we get finished with the other relationship started in this section.

   You can demonstrate that the relationship is again uni-directional using the rails console. We can navigate from the `Entrant` to the `Racer` but no longer have a way to navigate back.

   Note that this will still break lecture2_spec tests for rq06 since there is neither a races relationship or method at the moment. This will be repaired shortly.

   ```
   > entrant=Entrant.first
   > racer=entrant.racer
   > racer.methods.grep /races/
    => []

   $ rspec spec/lecture3_spec.rb -e rq03
   ```

4. Complete the demonstration of the 1:M embedded relationship between `Contest` and `Entrant`.

   Using the rails console, you can demonstrate your 1:M embedded relationship, with the embedded `Entrant` class having a link to the `Racer`. Notice how each of the `Entrants` are now stored within an array contained in the `Contest` document.

   ```
   > ["one","two"].each {|lname| Racer.create(:fn=>"thing",:ln=>lname) }
   > contest=Contest.first
   > contest.entrants.create(:id=>1, :group=>"youth", :racer=>Racer.find_by(:ln=>"one"))
   > contest.entrants.create(:id=>2, :group=>"youth", :racer=>Racer.find_by(:ln=>"two"))
   > contest.entrants.create(:id=>0, :group=>"masters", :racer=>Racer.find_by(:fn=>"cat"))
   > pp Contest.first.attributes
   {"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
    "name"=>"Boston 5K",
    "date"=>2015-05-30 00:00:00 UTC,
    "updated_at"=>2015-12-20 06:07:15 UTC,
    "venue_id"=>nil,
    "entrants"=>
     [{"_id"=>1,
       "group"=>"youth",
       "racer_id"=>BSON::ObjectId('5676438be301d0a1fb000026'),
       "name"=>"one, thing"},
      {"_id"=>2,
       "group"=>"youth",
       "racer_id"=>BSON::ObjectId('5676438be301d0a1fb000027'),
       "name"=>"two, thing"},
      {"_id"=>0,
       "group"=>"masters",
       "racer_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004'),
       "name"=>"inhat, cat"}]}

   $ rspec spec/lecture3_spec.rb -e rq04
   ```

5. Add an instance method called `races` that re-implements access to the now embedded `Entrant` documents from `Racer` using a query and application logic (as the error message suggests.) This method must:

   - be called `races`
   - query for all `Contest` instances that contain an `Entrant` with a foreign key (`racer_id`) equal to the primary key (`_id`) of the current `Racer` instance (**Hint**: `where` function) and
   - return references to only the embedded documents (**Hint**: `map` function)

   You can demonstrate your new navigation using the rails console. Given a `Racer`, you should be able to obtain a collection of `Entract` instances for races.

   ```
   > racer=Racer.find_by(:fn=>"cat",:ln=>"inhat")
   > racer.races
   ```

```
=> [#<Entrant _id: 0, name: "inhat, cat", group: "masters", secs: nil,
  racer_id: BSON::ObjectId('5675ff11e301d0a1fb000004')>]
> racer.races.first.contest.name
=> "Boston 5K"
```

We can verify that list by coming from the other direction.

```
> contest=Contest.where(:"entrants.name"=>{:$regex=>"inhat"}).first
> contest.entrants.where(:name=>"inhat, cat").first
=> #<Entrant _id: 0, name: "inhat, cat", group: "masters", secs: nil,
  racer_id: BSON::ObjectId('5675ff11e301d0a1fb000004')>
```

```
$ rspec spec/lecture3_spec.rb -e rq05
```

Lecture 4: 1:1 Linked

In this section we will concentrate on creating a 1:1 linked relationship. We will create a `MedicalRecord`, which we want to keep associated with the `Racer`, but separate from the `Racer`.

1. Use the `rails generate` command to create a model class called `MedicalRecord` that is mapped to the `medical` collection and uses the following fields:

   - `conditions` : Array

   You can demonstrate your new model class using the rails console.

   ```
   > MedicalRecord.new(:conditions=>["A","B"])
   => #<MedicalRecord _id: 567655c6e301d0a1fb00002b, conditions: ["A", "B"]>
   ```

   Notice too that the documents in this collection are being stored in an alternate-named collection called `medical` and not the default name of `medicalrecords` derived from the model class name.

   ```
   > MedicalRecord.collection.name
   => "medical"
   ```

   ```
   $ rspec spec/lecture4_spec.rb -e rq01
   ```

2. Add a 1:1 linked, uni-directional relationship from `MedicalRecord` to `Racer` using the `belongs_to` macro. This relationship must:

   - be called `racer`

   You can demonstrate your new relationship using the rails console. The `MedicalRecord` class was implemented to store the foreign key of the `Racer`.

   ```
   > m=MedicalRecord.create(:conditions=>["A","B"])
   => #<MedicalRecord _id: 5676561ce301d0a1fb00002c, conditions: ["A", "B"],
       racer_id: nil>
   > m.racer=Racer.find_by(:fn=>"cat",:ln=>"inhat")
   > m.save
   => true
   > pp MedicalRecord.find(m.id).attributes
   {"_id"=>BSON::ObjectId('5676561ce301d0a1fb00002c'),
    "conditions"=>["A", "B"],
    "racer_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004')}
   ```

   ```
   $ rspec spec/lecture4_spec.rb -e rq02
   ```

3. Implement the inverse side of the 1:1 linked relationship in the `Racer` model class using the `has_one` macro to make it bi-directional. This relationship must:

   - be called `medical_record`

   You can demonstrate your new bi-directional relationship capability using the rails console. In the following example, we locate the `MedicalRecord` created in the previous step, navigate to the `Racer` and back again to print the conditions.

```
> m=MedicalRecord.first
> m.racer.medical_record.conditions
 => ["A", "B"]
```

Note that if we get rid of the original relationship...

```
racer=Racer.find_by(:fn=>"cat",:ln=>"inhat")
racer.medical_record.destroy
```

...we can create a new related instance using the create_(relationship) macro.

```
> racer.create_medical_record(:conditions=>["C","D"])
> pp MedicalRecord.first.attributes
{"_id"=>BSON::ObjectId('56765840e301d0b210000000'),
 "conditions"=>["C", "D"],
 "racer_id"=>BSON::ObjectId('5675ff11e301d0a1fb000004')}
```

```
$ rspec spec/lecture4_spec.rb -e rq03
```

Lecture 5: M:M Linked

In this section we will focus on implementing a relationship where the remote foreign key is stored on both sides of the relationship as a part of an M:M bi-directional relationship.

1. Use the `rails generate` command to create a model class called `Judge` with the following fields:

   - `first_name` : String
   - `last_name` : String

   You can demonstrate your new model class using the rails console.

   ```
   > Judge.new
    => #<Judge _id: 56765acee301d0b210000001, first_name: nil, last_name: nil>
   ```

   ```
   $ rspec spec/lecture5_spec.rb -e rq01
   ```

2. Create a M:M uni-directional, linked relationship between `Judge` and `Contest` using the `has_and_belongs_to_many` macro. This relationship must:

   - be called `contests`

   You can demonstrate your new relationship using the rails console. Notice that when you form the relationship – the foreign key is stored in an array on what would have been the parent. This array implements the `belongs_to` part of `has_and_belongs_to_many`.

   ```
   > judge=Judge.create(:first_name=>"Judy")
   > contest=Contest.first
   > judge.contests << contest
   > pp Judge.first.attributes
   {"_id"=>BSON::ObjectId('56765bd0e301d0b210000003'),
    "first_name"=>"Judy",
    "contest_ids"=>[BSON::ObjectId('56761944e301d0a1fb00000d')]}
   ```

   ```
   $ rspec spec/lecture5_spec.rb -e rq02
   ```

3. Implement the reverse direction of the M:M linked relationship using the same `has_and_belongs_to_many` macro on the alternate side. The relationship must:

   - be called `judges`

   You can demonstrate your new navigation using the rails console. **Note** how Mongoid updates both sides of the relation except this time it is smart enough to know that the judge already has a record of the relationship.

   ```
   > judge=Judge.first
   > contest=Contest.first
   ```

```
> contest.judges << judge
> pp Contest.first.attributes
{"_id"=>BSON::ObjectId('56761944e301d0a1fb00000d'),
 ...
 "judge_ids"=>[BSON::ObjectId('56765bd0e301d0b210000003')]}
> pp Judge.first.attributes
{"_id"=>BSON::ObjectId('56765bd0e301d0b210000003'),
 "first_name"=>"Judy",
 "contest_ids"=>[BSON::ObjectId('56761944e301d0a1fb00000d')]}

$ rspec spec/lecture5_spec.rb -e rq03
```

Lecture 6: Validations and Constraints

In this section we will focus our attention on adding more definition to fields and relationships we have in place. For the `Entrant` to `Racer` relationship – the `Entrant` relied on both `first_name` and `last_name` being present in the `Racer` but there was nothing that made that a requirement. In that same relationship, there was also no requirement that the `Racer` exist when the `Entrant` was added.

1. In the `Racer` class, define a built-in validation that tests the presence of `first_name` and `last_name`.

   You can demonstrate your new validation using the rails console. Notice in the example below we can instantiate an invalid `Racer` and inspect the object for validation errors. However, the validation does not yet occur until specifically called.

   ```
   > r=Racer.new
   => #<Racer _id: 5676d02ee301d0b21000000b,
      first_name(fn): nil, last_name(ln): nil, date_of_birth(dob): nil>
   > r.errors.messages
   => {}
   > r.validate
   => false
   > r.errors.messages
   => {:first_name=>["can't be blank"], :last_name=>["can't be blank"]}
   ```

   The same sequence is followed if we attempt to instantiate a new, invalid `Racer` and then `save` it. The `save` fails with a quiet status response we can check and then poll for the errors.

   ```
   > r=Racer.new
   > r.errors.messages
        => {}
   > r.save
   => false
   > r.errors.messages
   => {:first_name=>["can't be blank"], :last_name=>["can't be blank"]}
   ```

   To get a more violent error response back from Mongoid, we can call the `save!` method instead, which reports the error using an exception.

   ```
   > r.save!
   Mongoid::Errors::Validations:
   message:
     Validation of Racer failed.
   summary:
     The following errors were found: First name can't be blank, Last
     name can't be blank
   resolution:
     Try persisting the document with valid data or remove the
     validations.
   ```

   ```
   $ rspec spec/lecture6_spec.rb -e rq01
   ```

14

2. Update the relationship from `Entrant` to `Racer` so that `Racer` will be validated each time `Entrant` is validated.

You can demonstrate your validation cascades across relationships using the rails console. In the following, we instantiate an `Entrant` with an invalid `Racer`. The `Entrant` is reported as an error because it references an invalid `Racer`. We can walk the relationship to the `Racer` to determine more details about its validation error.

```
> entrant=Entrant.create(:racer=>Racer.new)
> entrant.validate
 => false
> entrant.errors.messages
 => {:racer=>["is invalid"]}
> entrant.racer.errors.messages
 => {:first_name=>["can't be blank"], :last_name=>["can't be blank"]}

$ rspec spec/lecture6_spec.rb -e rq02
```

3. Update the relationship from `MedicalRecord` to `Racer` so that the relationship is required to exist for `MedicalRecord` to be valid.

You can demonstrate your required relationship using the rails console. In the example below, `MedicalRecord` is invalid because it does not have a relationship to a `Racer`.

```
> m=MedicalRecord.create
 => #<MedicalRecord _id: 5676d759e301d0b210000016, conditions: nil, racer_id: nil>
> m.validate
 => false
> m.errors.messages
 => {:racer=>["can't be blank"]}

$ rspec spec/lecture6_spec.rb -e rq03
```

4. Update the relationship from `Venue` to `Contest` to have `Venue` restricted from being deleted if there is a child `Contest` referencing it.

Your new relationship constraint can be demonstrated using the rails console. In the example below, we create a new `Venue` and associated `Contest`. `Contest` is stored with the foreign key to `Venue`.

```
> venue=Venue.create(:name=>"Flat Field")
 => #<Venue _id: 5676dd57e301d0b21000001c, name: "Flat Field">
> venue.contests.create(:name=>"FF10K")
 => #<Contest _id: 5676dd8ee301d0b21000001d,
    updated_at: 2015-12-20 16:55:42 UTC, name: "FF10K", date: nil,
    venue_id: BSON::ObjectId('5676dd57e301d0b21000001c'), judge_ids: nil>
```

Starting from a fresh find (so that we see what Mongoid needs to do to enforce the relationship) – we can observe that Mongoid performs a count in the database for all `Contest` documents containing the foreign key to this `Venue` instance. It will report an error when the count is greater than 0.

```
> reload!
> venue=Venue.find_by(:name=>"Flat Field")
D, | {"find"=>"venues", "filter"=>{"name"=>"Flat Field"}}

> venue.destroy
D, | {"count"=>"contests",
   "query"=>{"venue_id"=>BSON::ObjectId('5676dd57e301d0b21000001c')}}
Mongoid::Errors::DeleteRestriction:
message:
   Cannot delete Venue because of dependent 'contests'.
summary:
   When defining 'contests' with a :dependent => :restrict, Mongoid will
   raise an error when attempting to delete the Venue when the child
   'contests' still has documents in it.
resolution:
```

```
        Don't attempt to delete the parent Venue when it has children, or
        change the dependent option on the relation.
```

If we delete the relationship between the `Contest` and `Venue`, we see the foreign key being erased from the `Contest`.

```
> contest=venue.contests.first
> venue.contests.delete contest
 => #<Contest _id: 5676dd8ee301d0b21000001d, updated_at: 2015-12-20 16:57:39 UTC,
    name: "FF10K", date: nil, venue_id: nil, judge_ids: nil>
```

A re-attempt to delete the `Venue` causes the count query to return 0 and the delete of the `Venue` completes successfully.

```
> venue.destroy
D, | {"count"=>"contests", "query"=>{"venue_id"=>BSON::ObjectId('5676dd57e301d0b21000001c')}}
D, | {"delete"=>"venues", "deletes"=>[{"q"=>{"_id"=>BSON::ObjectId('5676dd57e301d0b21000001c')}...
 => true
```

```
$ rspec spec/lecture6_spec.rb -e rq04
```

5. Update the relationship from `Racer` to `MedicalRecord` to have the `MedicalRecord` destroyed when a `Racer` is removed.

   Your new relationship constraint can be demonstrated using the rails console. In the example below, we create a new `Racer` and associated `MedicalRecord`. The foreign key to the `Racer` is stored within the `MedicalRecord`.

```
> racer=Racer.create(:fn=>"Sally",:ln=>"Walden")
> racer.create_medical_record(:conditions=>["messy"])
```

   If we reload our state, re-find our `Racer`, and delete that instance, Mongoid will load the `MedicalRecord`, perform any callbacks on that `MedicalRecord`, and then delete both documents from the separate collections.

```
> reload!
> racer=Racer.where(:id=>racer.id).first
> racer.delete
D, | {"find"=>"medical", "filter"=>{"racer_id"=>BSON::ObjectId('5676e4dce301d0b210000027')}}
D, | {"delete"=>"medical", "deletes"=>[{"q"=>{"_id"=>BSON::ObjectId('5676e4e1e301d0b210000028')}...
D, | {"delete"=>"racers", "deletes"=>[{"q"=>{"_id"=>BSON::ObjectId('5676e4dce301d0b210000027')}...
 => true
```

```
$ rspec spec/lecture6_spec.rb -e rq05
```

Self Grading/Feedback

Unit tests have been provided in the bootstrap files that can be used to evaluate your solution. They must be run from the same directory as your solution.

```
$ rspec
........

(N) examples, 0 failures
```

Submission

There is no submission required for this assignment but the skills learned will be part of a follow-on assignment so please complete this to the requirements of the unit test.

Last Updated: 2015-01-16