


# Web-based State using Hidden Fields and Cookies

# HTTP is *stateless*

- The lifetime of an HTTP **session** is typically:
  - Client connects to the server
  - Client issues a request
  - Server responds
  - Client issues a request for something in the response
  - .... repeat ....
  - Client disconnects
- HTTP has no means of noting “oh this is the same client from that previous session”
  - *How is it you don't have to log in at every page load?*

# Maintaining State



- **Web application maintains *ephemeral* state**
  - Server processing often produces intermediate results
    - Not ACID, long-lived state
  - **Send** such **state to the client**
  - Client **returns the state** in subsequent **responses**


Two kinds of state: **hidden fields**, and **cookies**

# Ex: Online ordering

[socks.com/order.php](http://socks.com/order.php)



[socks.com/pay.php](http://socks.com/pay.php)



Separate page

# Ex: Online ordering

## What's presented to the user

pay.php

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="5.50">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

# Ex: Online ordering

## The corresponding backend processing

```
if(pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

# Ex: Online ordering

## What's presented to the user

```
<html>
<head> <title>Pay</title> </head>
<body>

<form action="submit_order" method="GET">
The total cost is $5.50. Confirm order?
<input type="hidden" name="price" value="0.01">
<input type="submit" name="pay" value="yes">
<input type="submit" name="pay" value="no">

</body>
</html>
```

Client can change  
the value!

# Solution: *Capabilities*

- **Server maintains *trusted state*** (while client maintains the rest)
  - Server stores intermediate state
  - Send a **capability** to access that state to the client
  - Client **references the capability** in subsequent responses
- **Capabilities should be large, random numbers**, so that they are hard to guess
  - To prevent illegal access to the state



# Using capabilities

## What's presented to the user

```
<html>
<head> <title>Pay</title> </head>
<body>
```

```
<form action="submit_order" method="GET">
```

The total cost is \$5.50. Confirm order?

```
<input type="hidden" name="sid" value="781234">
```

```
<input type="submit" name="pay" value="yes">
```

```
<input type="submit" name="pay" value="no">
```

```
</body>
```

```
</html>
```

**Capability;**  
the system will  
detect a change  
and abort

# Using capabilities


## The corresponding backend processing

```
price = lookup(sid);
if(pay == yes && price != NULL)
{
    bill_creditcard(price);
    deliver_socks();
}
else
    display_transaction_cancelled_page();
```

### **But: we don't want to pass hidden fields around all the time**

- Tedious to add/maintain on all the different pages
- Have to start all over on a return visit (after closing browser window)

# Statefulness with Cookies



- Server **maintains trusted state**
  - Server indexes/denotes state with a **cookie**
  - Sends cookie to the client, which stores it
  - Client returns it with subsequent queries to that same server

# Cookies are key-value pairs

Set-Cookie: **key**=**value**; **options**; ....

Headers


Data

```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqcali0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; path=/; domain=zdnet.com
Set-Cookie: user_agent=desktop
Set-Cookie: zdnet_ad_session=f
Set-Cookie: firstpg=0
Expires: Thu, 19 Nov 1981 08:52:00 GMT
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Pragma: no-cache
X-UA-Compatible: IE=edge,chrome=1
Vary: Accept-Encoding
Content-Encoding: gzip
Content-Length: 18922
Keep-Alive: timeout=70, max=146
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8
```

```
<html> ..... </html>
```

# Cookies

Set-Cookie: `edition=us`; `expires=Wed, 18-Feb-2015 08:20:34 GMT`; `path=/`; `domain=.zdnet.com`



## Semantics

- Store "us" under the key "edition"
- This value is no good as of Wed Feb 18...
- This value should only be readable by any domain ending in `.zdnet.com`
- This should be available to any resource within a subdirectory of /
- Send the cookie with any future requests to `<domain>/<path>`

# Requests with cookies

```
HTTP/1.1 200 OK
Date: Tue, 18 Feb 2014 08:20:34 GMT
Server: Apache
Set-Cookie: session-zdnet-production=6bhqca1i0cbciagu11sisac2p3; path=/; domain=zdnet.com
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0
Set-Cookie: edition=us; expires=Wed, 18-Feb-2015 08:20:34 GMT; path=/; domain=.zdnet.com
Set-Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11; path=/; domain=zdnet.com
```



**Subsequent visit**

## HTTP Headers

http://zdnet.com/

GET / HTTP/1.1

Host: zdnet.com

User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.9.2.11) Gecko/20101013 Ubuntu/9.04 (jaunty) Firefox/3.6.11

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,\*/\*;q=0.8

Accept-Language: en-us,en;q=0.5

Accept-Encoding: gzip,deflate

Accept-Charset: ISO-8859-1,utf-8;q=0.7,\*;q=0.7

Keep-Alive: 115

Connection: keep-alive

Cookie: session-zdnet-production=59ob97fpinqe4bg6lde4dvvq11 zdregion=MTI5LjluMTI5LjE1Mzp1czp1czpjZDJmNWY5YTdkODU1N2Q2YzM5NGU3M2Y1ZTRmN0 ...

# Why use cookies?

- **Session identifier**

- After a user has authenticated, subsequent actions provide a cookie
- So the user does not have to authenticate each time

- **Personalization**

- Let an anonymous user customize your site
- Store font choice, etc., in the cookie

# Why use cookies?

- **Tracking users**

- Advertisers want to know your behavior
- Ideally build a profile *across different websites*
  - Visit the Apple Store, then see iPad ads on Amazon?!
- How can site B know what you did on site A?

A shows you an ad from B and B scrapes the referrer URL

Option 1: B maintains a DB,  
indexed by your IP address

Option 2: B maintains a DB  
indexed by a *cookie*

**Problem: IP addrs change**

- **“Third-party cookie”**
- **Commonly used by large ad networks (doubleclick)**

<http://live.wsj.com/video/how-advertisers-use-internet-cookies-to-track-you>