

Symbolic execution systems

Resurgence

- Two key systems that triggered revival of this topic:
- **DART** — Godefroid and Sen, PLDI 2005
 - Godefroid = model checking, formal systems background
- **EXE** — Cadar, Ganesh, Pawlowski, Dill, and Engler, CCS 2006
 - Ganesh and Dill = SMT solver called STP (used in implementation), Cadar and Engler = systems
- Now on to next-generation systems

SAGE

- **Concolic executor** developed at **Microsoft Research**
 - Grew out of Godefroid's work on DART
 - Uses generational search
- Primarily **targets bugs in file parsers**
 - E.g., JPEG, DOCX, PPT, etc
 - Good fit for concolic execution
 - Likely to terminate
 - Just input/output behavior

SAGE Impact

- **Used on production software at MS.** Since 2007:
 - 500+ machine years (in largest fuzzing lab in the world)
 - Large cluster of machines continually running SAGE
 - 3.4 Billion+ constraints (largest SMT solver usage ever!)
 - 100s of apps, 100s of bugs (missed by everything else...)
 - Ex: *1/3 of all Win7 WEX security bugs found by SAGE*
 - Bug fixes shipped quietly to 1 Billion+ PCs
 - Millions of dollars saved (for Microsoft and the world)
 - SAGE is now used daily in Windows, Office, etc.

http://research.microsoft.com/en-us/um/people/pg/public_psfiles/SAGE-in-1slide-for-PLDI2013.pdf

KLEE

- **Symbolically executes LLVM bytecode**
 - LLVM compiles source file to .bc file
 - KLEE runs the .bc file
 - Grew out of work on EXE
- Works in the style of our basic symbolic executor
 - Uses `fork()` to manage multiple states
 - Employs a variety of search strategies
 - Primarily **random path + coverage-guided**
 - Mocks up the environment to deal with system calls, file accesses, etc.
- **Freely available with LLVM distribution**

KLEE: Coverage for Coreutils




Figure 6: Relative coverage difference between KLEE and the COREUTILS manual test suite, computed by subtracting the executable lines of code covered by manual tests (L_{man}) from KLEE tests (L_{klee}) and dividing by the total possible: $(L_{klee} - L_{man})/L_{total}$. Higher bars are better for KLEE, which beats manual testing on all but 9 applications, often significantly.

Cadar, Dunbar, and Engler. KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs, OSDI 2008

KLEE: Coreutils crashes

<code>paste -d\\ abcdefghijklmnopqrstuvwxyz</code>
<code>pr -e t2.txt</code>
<code>tac -r t3.txt t3.txt</code>
<code>mkdir -Z a b</code>
<code>mkfifo -Z a b</code>
<code>mknod -Z a b p</code>
<code>md5sum -c t1.txt</code>
<code>ptx -F\\ abcdefghijklmnopqrstuvwxyz</code>
<code>ptx x t4.txt</code>
<code>seq -f %0 1</code>

<code>t1.txt: "\t \tMD5 ("</code>
<code>t2.txt: "\b\b\b\b\b\b\b\b\t"</code>
<code>t3.txt: "\n"</code>
<code>t4.txt: "a"</code>

Figure 7: KLEE-generated command lines and inputs (modified for readability) that cause program crashes in COREUTILS version 6.10 when run on Fedora Core 7 with SELinux on a Pentium machine.

Mayhem

- Developed at CMU (Brumley et al), **runs on binaries**
- Uses BFS-style search and native execution
 - **Combines best of symbolic and concolic strategies**
- **Automatically generates exploits** when bugs found

Mergepoint

- Extends Mayhem with a technique called **veritesting**
 - **Combines symbolic execution** with **static analysis**
 - Use static analysis for complete code blocks
 - Use symbolic execution for hard-to-analyze parts
 - Loops (how many times will it run?), complex pointer arithmetic, system calls
- Better **balance** of time **between solver and executor**
 - **Finds bugs faster**
 - **Covers more of the program** in the same time
- Found 11,687 bugs in 4,379 distinct applications in a Linux distribution
 - Including new bugs in highly tested code

Other symbolic executors

- **Cloud9** — Parallel, multi-threaded symbolic execution
 - Extends KLEE (available)
- **jCUTE, Java PathFinder** — symbolic execution for Java (available)
- **Bitblaze** — Binary analysis framework (available)
- **Otter** — directed symbolic execution for C (available)
 - Give the tool a line number, and it try to generate a test case to get there
- **Pex** — symbolic execution for .NET

Summary

- **Symbolic execution generalizes testing**
 - Uses static analysis to direct generation of tests that cover different program paths
- Used in practice to find **security-critical bugs** in **production code**
 - SAGE at Microsoft
 - Mergepoint for Linux
- **Many tools freely available**